

CS395T: Introduction to Scientific and Technical Computing

Instructors:

Dr. Karl W. Schulz, Research Associate, TACC

Dr. Victor Eijkhout, Research Scientist, TACC

Outline

- Finish Architecture Discussion
- Ranger Intro - *pictures O' the week*
- Class Accounts
 - *Do's and Don'ts*
- Batch System Introduction and Usage

Data Reuse

- Performance is limited by data transfer rate
- High performance if data items are used multiple times
- Example: vector addition $x_i = x_i + y_i$: 1 op, 3 mem accesses
- Example: inner product $s = s + x_i * y_i$: 2 op, 2 mem access (s in register; also no writes)

Programming Strategies: Contiguous Access

- Avoid strides: cache lines contain 4 (or so) words, might as well use them all
- Example: dot product, sequential access of the vectors
- Strided dot product:

```
sum=0.;
```


```
for (j=0; j < stride; ++j)
```

```
    for(i=j; i < n; i+=stride)
```

```
        sum += a[i]*b[i];
```

Not all elements on a cache line used.

Dot Product Performance




What's Going On?

- Small vectors are noisy
 - probably not enough work to be measuring well
 - even still non-stride-1 access foil the plans of the hardware prefetcher
- Eventually everyone gets to the peak
 - ~1.8 GFlops = ~20% of peak
 - not far from what we predicted
 - probably some improvement yet to be had
- For *stride* $\neq 1$ we see the L1 (32K) cache size boundary
- For *stride* $= 1$ prefetching and other latency hiding tricks let the processor maintain performance
- Everybody hits the L2 (4MB) cache size boundary pretty hard

Programming Strategies: Blocked Algorithms

- Long vectors are flushed from cache: break up in smaller blocks
- Reuse of input vector (limited)
- Use of cache lines in matrix



More on blocked algorithms

- This gets tricky fast
- Matrix-matrix multiply is triple loop; blocked is 6-deep loop
- Choice of blocking sizes is complicated
- Loop exchange to aim for L1 or L2 reuse (Atlas vs Goto approach)
- [The Solution](#): Use optimized math libraries whenever possible to do routine linear algebra in your applications (more on libraries coming later in the class)

Pipelining

Pipeline

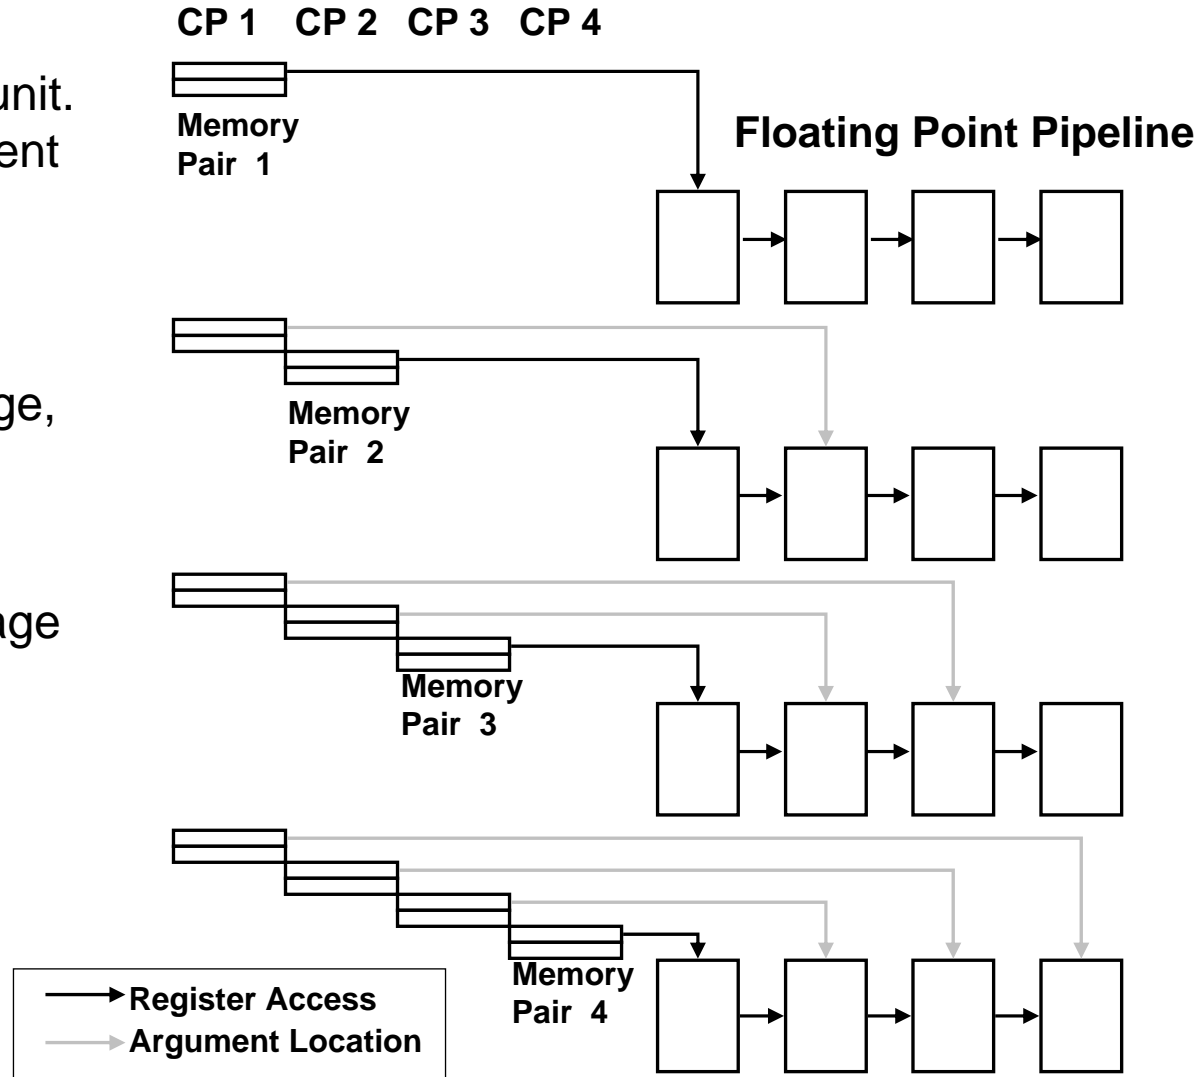
A serial multistage functional unit. Each stage can work on different sets of independent operands simultaneously.

After execution in the final stage, first result is available.

Latency = # of stages * CP/stage

CP/stage is the same for each stage and usually 1.

4-Stage FP Pipe



Branch Prediction

- The “instruction pipeline” is all of the processing steps (also called segments) that an instruction must pass through to be “executed”.
- Higher frequency machines have a larger number of segments.
- Branches are points in the instruction stream where the execution may jump to another location, instead of executing the next instruction.
- For repeated branch points (within loops), instead of waiting for the loop to branch route outcome, it is predicted.

Pentium III processor pipeline

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|----|
| 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 |
|---|---|---|---|---|---|---|---|---|----|

Pentium 4 processor pipeline


| | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|----|----|----|----|----|----|----|----|----|----|----|
| 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 | 15 | 16 | 17 | 18 | 19 | 20 |
|---|---|---|---|---|---|---|---|---|----|----|----|----|----|----|----|----|----|----|----|

Misprediction is more “expensive” on Pentium 4’s.

TACC's Ranger System

First NSF Track2 System: 500+ Tflops!

- TACC selected for first NSF 'Track2' HPC system
 - \$30M system acquisition
 - Sun is the vendor
 - Competed against almost every open science HPC center
- TACC, ICES, Cornell, ASU operating/supporting system four 4 years (\$29M)



Ranger System Configuration

- **Compute Power** - 504 Teraflops Peak Performance
 - 3,936 Sun four-socket, quad-core compute nodes
 - 15,744 AMD Opteron “Barcelona” processors
 - Quad-core, four flops/cycle (dual pipelines)
- **Memory**
 - 2 GB/core, 32 GB/node, 125 TB total
 - 132 GB/s aggregate bandwidth
- **Infiniband interconnect**
 - Full non-blocking 7-stage Clos fabric
 - Low latency ($\sim 2.3 \mu\text{sec}$), high-bandwidth ($\sim 950 \text{ MB/s}$)

Impact in NSF TeraGrid

- 460M CPU hours to TeraGrid per year
 - more than double current total capacity of all TG HPC systems
 - 1.8 Billion CPU hours over operational life
- 529 Teraflops peak
 - 2x total performance of all TG HPC systems
 - 8x top TG HPC system in performance, memory, disk
- Balanced, general-purpose capability system
 - More than 60,000 cores available
 - Unprecedented scaling opportunities for computational science and research
- Re-establish NSF as a leader in HPC
- *Jumpstarts progress to petascale for entire US academic research community*

Ranger User Environment

- The overall look and feel of *Ranger* from the user perspective will be very similar to our current Linux cluster
 - Full Linux OS w/ hardware counter patches on login and compute nodes (2.6.12.6 is starting working kernel)
 - Lustre File System
 - \$HOME, and multiple \$WORKS will be available
 - Largest \$WORK will be ~1PB total
 - Standard 3rd party packages
 - Infiniband using next generation of Open Fabrics
 - MVAPICH and OpenMPI (MPI1 and MPI2)
- Suite of compilers
 - Portland Group PGI
 - Sun Studio
 - PathScale
 - *Possibly the Intel compiler*

Ranger Disk Subsystem - *Lustre*

- Disk system (OSS) is based on Sun x4500 “Thumper” servers - similar to TiTech installation
 - Each server has 48 SATA II 500 GB drives (24TB total) - running internal software RAID
 - Dual Socket/Dual-Core Opterons @ 2.6 GHz
 - Downside is that these nodes have PCI-X - raw I/O bandwidth can exceed a single PCI-X 4X Infiniband HCA
 - 72 Servers Total: 1.7 PB raw storage
- Metadata Servers (MDS) based on Sun Fire x4600s
- MDS is Fibre-channel connected to 9TB Flexline Storage
- Target Performance
 - Aggregate bandwidth: 40 GB/sec



Design:

- Top loading Disks
- Front to rear airflow
- Redundant fans
- Passive Backplane
- No wires in box

Reliability/Availability

- Enterprise class SATA disks
- 1M hours MTBF
- RAID 0, 1, 5, 10
- Redundant Power
- Hot-swap FRUs

Ranger System Configuration

| Logical Volume Name | Estimated Raw Capacity | Target Usage |
|---------------------|------------------------|---|
| <i>WORK1</i> | ~850 PB | Large temporary storage; not backed up, purged periodically |
| <i>WORK2</i> | ~250 TB | Large allocated storage; not backed up, quota enforced |
| <i>PROJECTS</i> | 2 TB | Repository for TeraGrid Community Software |
| <i>HOME1</i> | 2 TB | Permanent user storage; automatically backed up, quota enforced |
| <i>HOME2</i> | 2 TB | Permanent user storage; automatically backed up, quota enforced |
| <i>HOME3</i> | 2 TB | Permanent user storage; automatically backed up, quota enforced |

Ranger Space, Power and Cooling

- System Power: 3.0 MW total
- System: 2.4 MW
 - ~90 racks, in 6 row arrangement
 - ~100 in-row cooling units
 - ~4000 ft² total footprint
- Cooling: ~0.6 MW
 - In-row units fed by three 400-ton chillers
 - Enclosed hot-aisles
 - Supplemental 280-tons of cooling from CRAC units
- Observations:
 - Space less an issue than power
 - Cooling > 25kW per rack difficult
 - Power distribution a challenge, more than 1200 circuits

Ranger Project Timeline

| | |
|-------|------------------------------|
| Sep06 | award, press, relief, beers |
| 1Q07 | equipment begins arriving |
| 2Q07 | facilities upgrades complete |
| 3Q07 | very friendly users |
| 4Q07 | more early users |
| Dec07 | production, many beers |
| Jan08 | allocations begin |

Note: all US academics are eligible to apply for a TeraGrid/Ranger allocation:

<https://pops-submit.ci-partnership.org/>

Ranger: External Infrastructure



Ranger: PDUs and In-Row Coolers



- APC In-row coolers are installed and plumbed
- Compute Racks will slide in between the coolers which are heat exchangers drawing from the hot aisles, exhausting ambient into the cold aisles

Ranger: Machine Room Layout



Computer Accounts

Computer Accounts

- Lonestar
 - Production Resource for UT and the NSF TeraGrid Community
 - 1460 Dell 1955 nodes
 - 2 dual-core 2.6GHz Intel Xeon (Woodcrest) Processors/node
 - 8 GB RAM/node
 - Infiniband Interconnect
 - Currently listed as #15 on the Top500 List

“Production”

- Jobs run in a managed environment
 - login to the login node
 - submit jobs to the scheduler
 - wait
 - collect results
- Running programs on the login node highly discouraged
 - avoid resource intensive tasks
 - exceptions include compilers, “standard” UNIX commands (ls, mkdir, cp, mv, etc.)

Lonestar Login

- SSH only
- UNIX users
 - ssh username@lonestar.tacc.utexas.edu*
- Windows users: Get a client
 - PuTTY
 - <http://www.chiark.greenend.org.uk/~sgtatham/putty/>
 - PuTTY can be gotten from Bevoware as well
 - <https://www.utexas.edu/its/bevoware/download/>
 - or use Cygwin and follow the UNIX instructions
 - <http://www.cygwin.com/>

First Login

- Login using your username and password
 - both are **case sensitive**

- Change your password

```
lslogin1% passwd
```

```
Changing password for user istc00.
```

```
Changing password for istc00
```

```
(current) UNIX password:
```

```
New UNIX password:
```

```
Retype new UNIX password:
```

```
lslogin1%
```

- Logout

```
lslogin1% logout
```

Lonestar Usage

- If you don't know what you're doing, **WAIT!**
- If you're already running on Lonestar, feel free to use either your account or your class account for assignments
- **DON'T** use the class account to do your research
 - your advisor needs to setup a project and add you to it for research usage

Login Problems

- Send me email
 - karl@tacc.utexas.edu
 - eijkhout@tacc.utexas.edu
- Subject
 - CS395T Can't login to Lonestar
- Include any error message that you saw

Batch Systems

Batch Systems

- In a number of scientific computing environments, multiple users must *share* a compute resource:
 - research clusters
 - supercomputing centers
- On multi-user HPC clusters, the *batch system* is a key component for aggregating compute nodes into a single, sharable computing resource
- The batch system becomes the “nerve center” for coordinating the use of resources and controlling the state of the system in a way that must be “fair” to its users
- As current and future *expert* users of large-scale compute resources, you need to be familiar with the basics of a batch system

Batch Systems

- The core functionality of all batch systems are essentially the same, regardless of the size or specific configuration of the compute hardware:
 - Multiple Job Queues:
 - queues provide an orderly environment for managing a large number of jobs
 - queues are defined with a variety of limits for maximum run times, memory usage, and processor counts; they are often assigned different priority levels as well
 - may be interactive or non-interactive
 - Job Control:
 - submission of individual jobs to do some work (eg. serial, or parallel HPC applications)
 - simple monitoring and manipulation of individual jobs, and collection of resource usage statistics (e.g., memory usage, CPU usage, and elapsed wall-clock time per job)
 - Job Scheduling
 - policy which decides priority between individual user jobs
 - allocates resources to scheduled jobs

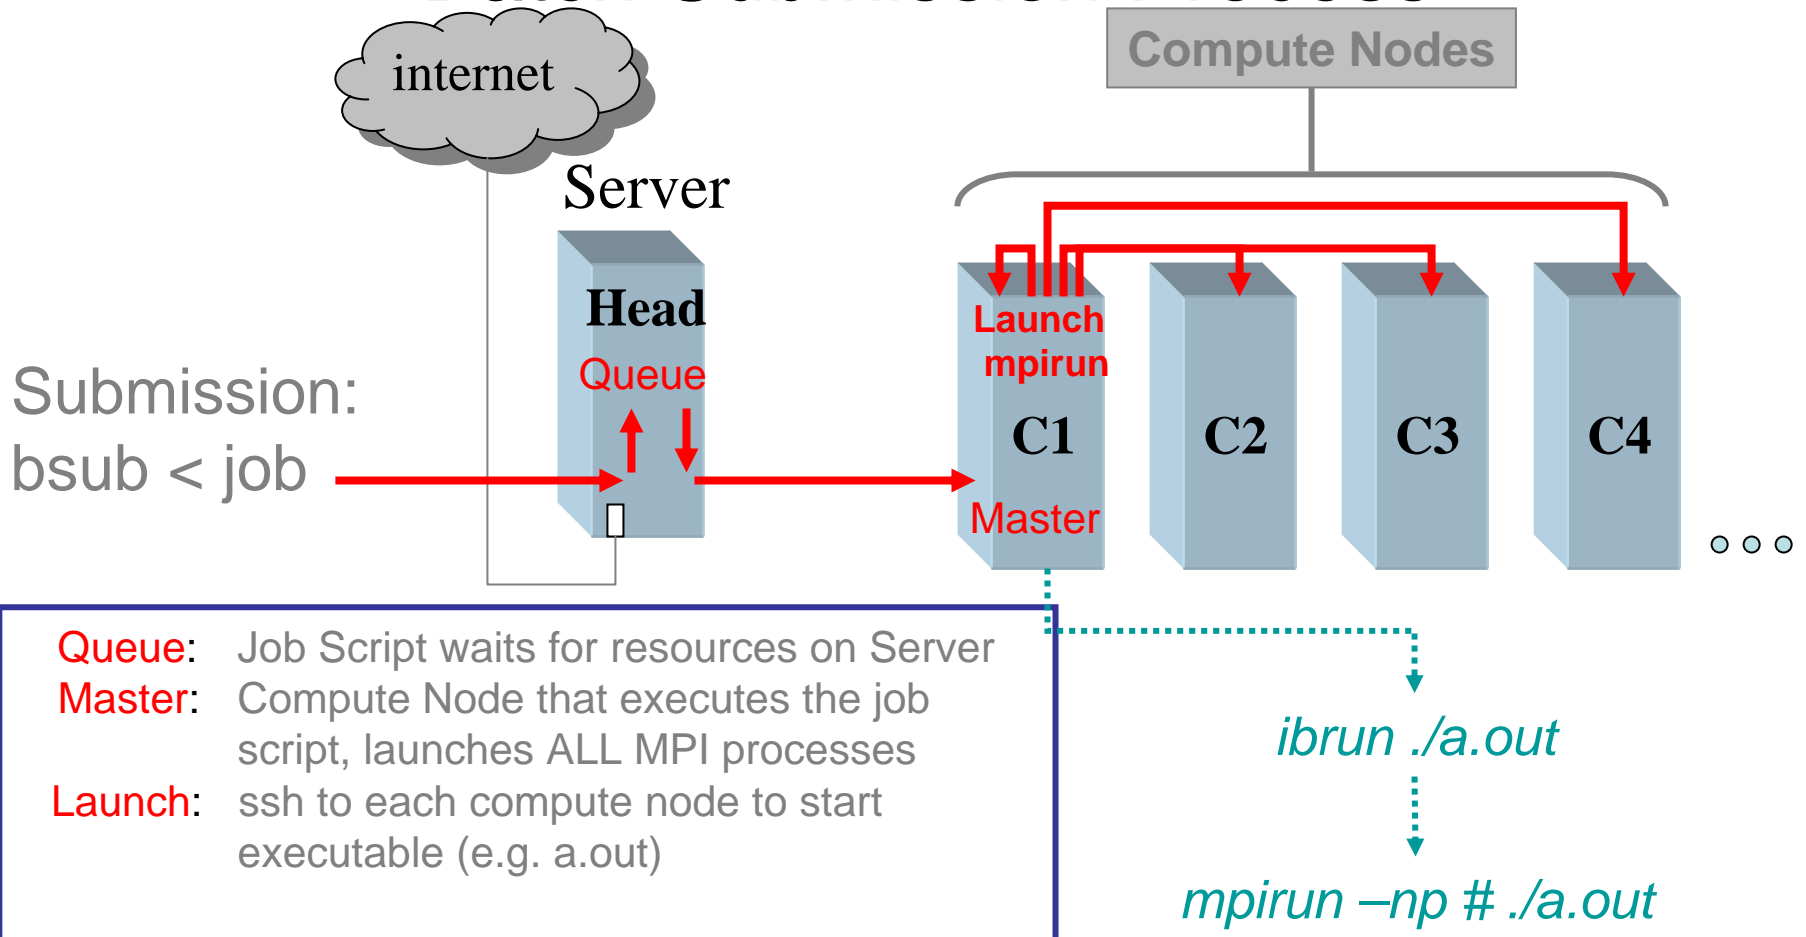
Batch Systems

- Job Scheduling Policies:
 - the scheduler must decide how to prioritize all the jobs on the system and allocate necessary resources for each job (processors, memory, file-systems, etc)
 - scheduling process can be easy or non-trivial depending on the size and desired functionality
 - **first in, first out (FIFO)** scheduling: jobs are simply scheduled in the order in which they are submitted
 - **political** scheduling: enables some users to have more priority than others
 - **fairshare** scheduling, scheduler ensures users have equal access over time
 - Additional features may also impact scheduling order:
 - **advanced reservations** - resources can be reserved in advance for a particular user or job
 - **backfill** - can be combined with any of the scheduling paradigms to allow smaller jobs to run while waiting for enough resources to become available for larger jobs
 - back-fill of smaller jobs helps maximize the overall resource utilization
 - back-fill can be your friend for small duration jobs

Batch Systems

- Common batch systems you may encounter in scientific computing:
 - Platform LSF
 - PBS
 - Loadleveler (IBM)
 - SGE
- All have similar functionality but different syntax
- Reasonably straight forward to convert your job scripts from one system to another
- Above all include specific batch system directives which can be placed in a shell script to request certain resources (processors, queues, etc).
- We will focus on LSF primarily since it is the system running on Lonestar

Batch Submission Process



LSF Batch System

- *Lonestar* uses Platform LSF for both the batch queuing system and scheduling mechanism (provides similar functionality to PBS, but requires different commands for job submission and monitoring)
- LSF includes global fairshare, a mechanism for ensuring no one user monopolizes the computing resources
- Batch jobs are submitted on the front end and are subsequently executed on compute nodes as resources become available
- Order of job execution depends on a variety of parameters:
 - Submission Time
 - Queue Priority: some queues have higher priorities than others
 - Backfill Opportunities: small jobs may be back-filled while waiting for bigger jobs to complete
 - Fairshare Priority: users who have recently used a lot of compute resources will have a lower priority than those who are submitting new jobs
 - Advanced Reservations: jobs may be blocked in order to accommodate advanced reservations (for example, during maintenance windows)
 - Number of Actively Scheduled Jobs: there are limits on the maximum number of concurrent processors used by each user

Lonestar Queue Definitions

| Queue Name | Max Runtime | Min/Max Procs | SU Charge Rate | Use |
|-------------|-------------|---------------|----------------|--|
| normal | 24 hours | 2/512 | 1.0 | Normal usage |
| high | 24 hours | 2/512 | 1.8 | Higher priority usage |
| development | 30 min | 1/32 | 1.0 | Debugging and development Allows <i>interactive</i> jobs |
| hero | 24 hours | >512 | 1.0 | Large job submission Requires special permission |
| serial | 12 hours | 1/1 | 1.0 | For serial jobs. No more than 4 jobs/user |
| request | | | | Special Requests |
| spruce | | | | Debugging & development, special priority, urgent comp. env. |
| systest | | | | System Use (<i>TACC Staff only</i>) |

LSF Fairshare

- A global fairshare mechanism is implemented on Lonestar to provide fair access to its substantial compute resources
- Fairshare computes a dynamic priority for each user and uses this priority in making scheduling decisions
- Dynamic priority is based on the following criteria
 - Number of shares assigned
 - Resources used by jobs belonging to the user:
 - Number of job slots reserved
 - Run time of running jobs
 - Cumulative actual CPU time (not normalized), adjusted so that recently used CPU time is weighted more heavily than CPU time used in the distant past

LSF Fairshare

- **bhpart:** Command to see current fairshare priority. For example:

```
lsl login1--> bhpart -r
```

```
HOST_PARTITON_NAME: Global Parti ti on
```

```
HOSTS: all
```

```
SHARE_I NFO_FOR: Gl obal Parti ti on/
```

Priority ↑

| USER/GROUP | SHARES | PRIORIT Y | STARTED | RESERVED | CPU_TI ME | RUN_TI ME |
|------------|--------|-----------|---------|----------|------------|-----------|
| avij i t | 1 | 0. 333 | 0 | 0 | 0. 0 | 0 |
| chona | 1 | 0. 333 | 0 | 0 | 0. 0 | 0 |
| ewal ker | 1 | 0. 333 | 0 | 0 | 0. 0 | 0 |
| mi nyard | 1 | 0. 333 | 0 | 0 | 0. 0 | 0 |
| phaa406 | 1 | 0. 333 | 0 | 0 | 0. 0 | 0 |
| bbarth | 1 | 0. 333 | 0 | 0 | 0. 0 | 0 |
| mi l fel d | 1 | 0. 333 | 0 | 0 | 2. 9 | 0 |
| karl | 1 | 0. 077 | 0 | 0 | 51203. 4 | 0 |
| vmcal o | 1 | 0. 000 | 320 | 0 | 2816754. 8 | 7194752 |

Commonly Used LSF Commands

| | |
|------------------------|---|
| bhosts | Displays configured compute nodes and their static and dynamic resources (including job slot limits) |
| lsload | Displays dynamic load information for compute nodes (avg CPU usage, memory usage, available /tmp space) |
| <bbsub< b=""></bbsub<> | submits a batch job to LSF |
| bqueues | displays information about available queues |
| bjobs | displays information about running and queued jobs |
| bhist | displays historical information about jobs |
| bstop | suspends unfinished jobs |
| brresume | resumes one or more suspended jobs |
| bkill | Sends signal to kill, suspend, or resume unfinished jobs |
| bhpart | Displays global fairshare priority |
| lshosts | Displays hosts and their static resource configuration |
| lsuser | Shows user job information |

*Note: most of these commands support a “-l” argument for long listings. For example: **bhist -l <job ID>** will give a detailed history of a specific job. Consult the man pages for each of these commands for more information.*

LSF Batch System

- LSF Defined Environment Variables:

| | |
|-----------------|---|
| LSB_ERRORFILE | name of the error file |
| LSB_JOBID | batch job id |
| LS_JOBPID | process id of the job |
| LSB_HOSTS | list of hosts assigned to the job. Multi-cpu hosts will appear more than once (may get truncated) |
| LSB_QUEUE | batch queue to which job was submitted |
| LSB_JOBNAME | name user assigned to the job |
| LS_SUBCWD | directory of submission, i.e. this variable is set equal to \$cwd when the job is submitted |
| LSB_INTERACTIVE | set to 'y' when the -I option is used with bsub |

LSF Batch System

- Comparison of LSF, PBS and Loadleveler commands that provide similar functionality

| LSF | PBS | Loadleveler |
|----------------------|---------------------------------------|------------------------|
| <code>bresume</code> | <code>qrls</code> <code>qsit</code> | <code>llhold -r</code> |
| <code>bsub</code> | <code>qsub</code> | <code>llsubmit</code> |
| <code>bqueues</code> | <code>qstat</code> | <code>llclass</code> |
| <code>bjobs</code> | <code>qstat</code> | <code>llq</code> |
| <code>bstop</code> | <code>qhold</code> | <code>llhold</code> |
| <code>bkill</code> | <code>qdel</code> | <code>llcancel</code> |

Batch System Concerns

- Submission (need to know)
 - Required Resources
 - Run-time Environment
 - Directory of Submission
 - Directory of Execution
 - Files for stdout/stderr Return
 - Email Notification
- Job Monitoring
- Job Deletion
 - Queued Jobs
 - Running Jobs

LSF: Basic MPI Job Script

| | | | |
|------------------------------|---|--------|--------------------------------------|
| <code>#!/bin/csh</code> | } |→ | Total number of processes |
| <code>#BSUB -n 32</code> | } |→ | Job name |
| <code>#BSUB -J hello</code> | } |→ | Stdout Output file name (%J = jobId) |
| <code>#BSUB -o %J.out</code> | } |→ | Stderr Output file name |
| <code>#BSUB -e %J.err</code> | } |→ | Submission queue |
| <code>#BSUB -q normal</code> | } |→ | Your Project Name |
| <code>#BSUB -P A-ccsc</code> | } |→ | Max Run Time (15 minutes) |
| <code>#BSUB -W 0:15</code> | } | | |

| | | |
|---|---|------------------------------------|
| <code>echo "Master Host = "`hostname`</code> | } | Echo pertinent environment info |
| <code>echo "LSF_SUBMIT_DIR: \$LS_SUBCWD"</code> | | |
| <code>echo "PWD_DIR: "`pwd`</code> | | |

| | | |
|----------------------------|---|-------------------|
| <code>ibrun ./hello</code> | } | Execution command |
|----------------------------|---|-------------------|

Parallel application manager and mpirun wrapper script

executable

LSF: Extended MPI Job Script

| | | | |
|--|---|---|---|
| <code>#!/bin/csh</code> | } | → | Total number of processes |
| <code>#BSUB -n 32</code> | } | → | Job name |
| <code>#BSUB -J hello</code> | } | → | Stdout Output file name (%J = jobId) |
| <code>#BSUB -o %J.out</code> | } | → | Stderr Output file name |
| <code>#BSUB -e %J.err</code> | } | → | Submission queue |
| <code>#BSUB -q normal</code> | } | → | Your Project Name |
| <code>#BSUB -P A-ccsc</code> | } | → | Max Run Time (15 minutes) |
| <code>#BSUB -W 0:15</code> | } | → | <i>Dependency on Job <1123></i> |
| <code>#BSUB -w 'ended(1123)'</code> | } | → | <i>Email address</i> |
| <code>#BSUB -u karl@tacc.utexas.edu</code> | } | → | <i>Email when job begins execution</i> |
| <code>#BSUB -B</code> | } | → | <i>Email job report information upon completion</i> |
| <code>#BSUB -N</code> | } | → | |

```
echo "Master Host = "`hostname`  
echo "LSF_SUBMIT_DIR: $LS_SUBCWD"
```

```
ibrun ./hello
```