

CS395T: Introduction to Scientific and Technical Computing

CVS

Instructors

Dr. Karl W. Schulz, TACC

Dr. Victor Eijkhout, TACC

Outline

- HW #1 Update
 - Expect to return on Thursday
- HW#2 is now Posted
 - You will need to login to Lonestar to complete
 - Submit results electronically on Blackboard
 - Due in 1 Week, 10/9/2007
- Source Code Control

Ranger Update



Source Control and Versioning

- Codes evolve over time
 - sometimes bugs creep in
 - sometimes the old way was right
 - sometimes it's nice to look back at the evolution
- How can you get back to an old version?
 - keep a copy of every version of every file
 - disk is cheap, but this could get out of hand quickly
 - is a huge pain to maintain
 - use a tool

Some Tools

- Free
 - RCS – Revision Control System
 - CVS – Concurrent Versions System
 - SVN – Subversion
 - ...
- Commercial
 - MS Visual SourceSafe
 - ClearCase
 - MKS
 - ...

Two Models of Source Control

- A central repository holds the files in both models
- Read/Write Local Workspaces and Merging
 - every developer has a local copy of the source files
 - everybody can read and write files in their local copy
 - conflicts between simultaneous edits handled with sophisticated merging algorithms or manually when files are synched against the repository or committed to it
- Read-only Local Workspaces and Locks
 - every developer has a read-only local copy of the source files
 - individual files are checked-out as needed and locked in the repo in order to gain write access
 - unlocking the file commits the changes to the repo and makes the file read-only again

CVS

- Started with some shell scripts in 1986
- Recoded in '89
- Evolving ever since (though mostly unchanging now)
- Uses r/w local workspaces and merging
- Only stores differences between versions
 - saves space
 - basically uses *diff(1)* and *diff3(1)*
- Works with local repositories or over the net with rsh/ssh

CVS and the Environment

- `$CVSROOT` gives the path to the repository
 - `/path/to/repo` for local repositories
 - `username@host:/path/to/repo` for remote
- `$CVS_RSH` tells CVS which remote shell command to use to talk to the remote machine
 - should almost always be set to `ssh`
- Both of these should generally be set before using `cv`s

Repositories

- CVS stores files in a central location called a repository
- You can create a repository anywhere you want
- Generally, this should be
 - accessible to your collaborators
 - backed up daily
 - sufficiently large
- You don't usually need more than one repo
- If you're lucky, your site or IT staff will provide a CVS repository already
- Repos are collections of modules, and modules contain your source code

Invoking CVS

```
cvcs [cvs-options] command [command-options-and-arguments]
```

- The first set of options are common to all *cvcs* commands
- The second set are command specific
- Some commands
 - init
 - checkout
 - update
 - commit
 - diff
 - add
 - log
 - tag
 - status
- Commands are usually recursive
 - current dir or its immediate subdirs must be CVS local working directories, or CVS will stop recursing

Creating a Repository

- Set `$CVSROOT` and `$CVS_RSH`

```
lslogin1$ cvs init
```

- You can also specify the repository on the command line

```
lslogin1$ cvs -d  
istc00@lonestar.tacc.utexas.edu: /home  
/utexas/istc/istc00/cvs import
```

Preparing to Import a New Module

- Clean up your source directory
- CVS isn't really designed for object files, executables, binary files, etc.
 - though it can deal with them if it has to
- It's designed for text files
 - source
 - scripts
 - *Makefiles*
 - text documents
 - etc.

Importing

```
lslogin1$ cd mydir
```

```
lslogin1$ cvs import myrepodir vendor_tag release_tag
```

- Creates a directory in *\$CVSROOT/myrepodir*
- Pulls in all of the files in *mydir*
- Applies two tags that can be used to get back to this exact version later
- Does not modify anything in *mydir*
 - Note: you have to check out a CVSified working copy to begin working on this source under CVS control

Checking Out a Module

```
lslogin1$ cd ..
```

```
lslogin1$ cvs checkout myrepodir
```

- Creates a directory called *myrepodir*
- Pulls down a copy of all the files in the repository in the *myrepodir* module
- Adds a directory for CVS's use called *myrepodir/CVS*
 - no reason to modify these files in most cases
 - ... but there are occasions when you will want to do so

```
lslogin1$ cvs co myrepodir
```

- Most CVS commands have a shorter name 'co' is short for '*checkout*'

Ignoring Certain Files

- In any given module, for each directory in that module, *.cvsignore* (notice the dot in the name) lists files and glob(7) patterns of files to be ignored
 - one file or pattern per line
 - every directory should get one that ignores tilde files, object files, and a.out

```
.cvsignore:
```

```
*~
```

```
.*~
```

```
*.o
```

```
a.out
```

Updating to the Latest

- Usually, before adding anything, or checking in your changes, you should synchronize the state of your local copy with things that may have changed in the repository version

```
lslogin1$ cvs update
```

```
lslogin1$ cvs up
```


Update

- Downloads new files
 - use *cv*s *up -d* to get new directories as well
- Downloads any differences in files between repo and the local copy
 - merges changes where possible
 - marks files that can't be merged
 - these need manual conflict resolution
- Marks locally modified files
- Marks locally added files
- Lists unknown files and directories
 - these probably need to be added or listed in *.cvsignore*

Update Codes

- U FILE
 - FILE was updated, whole file was downloaded
- P FILE
 - FILE was updated, only a “patch” was downloaded
- A FILE
 - FILE locally added with *cv*s *add*
 - this change needs to be committed
- M FILE
 - FILE has been locally modified
 - differences may have been successfully merged
 - if so, a backup of your local file is made
 - a message is printed telling you about it and giving the name of the backup

Update Codes

- C FILE
 - FILE has conflicts that require human intervention
 - CVS makes a backup of your original: `.#FILE.REVISION`
- ? FILE
 - CVS doesn't know anything about this file
 - add it
 - add it to `.cvsignore`
- R FILE
 - FILE has been removed with `cv remove FILE`
 - this change needs to be committed

Adding Files

```
lslogin1$ cvs add somefile
```

- Tells CVS to start tracking this file
- Does not actually upload to the repo
 - *somefile* is simply marked for addition to the repo
 - *cvs commit* needed to actually upload the file

Committing Changes

- Local modifications, added files, and removed files have to be committed to the repo

```
lslogin1$ cvs commit
```

```
lslogin1$ cvs ci
```

- Tries to update first
- If it finds conflicts, it will error and ask you to fix the conflicts before trying to continue

Commit

- CVS uses the *\$EDITOR* environment variable to invoke your favorite editor to allow you to enter a log message
 - if *\$EDITOR* is empty, it invokes *vi(1)*
 - if you don't like *vi(1)*, you should set this variable to something else

```
lslogin1$ export EDITOR=emacs
```

- You can specify the log message on the command line with the *-m* option to *commit*

```
lslogin1$ cvs ci -m "some message"
```

Viewing the Log Messages

```
lslogin1$ cvs log somefile
```

- Shows you
 - all the past log messages
 - who made the commit
 - version numbers
 - total revisions
 - etc.
- Doesn't change anything

```
[root@lonestar2 config]# cvs log config.tacc.sh
karl@cvs.tacc.utexas.edu's password:

RCS file: /cvs/adv/system/admin/lustre/lonestar2/config.tacc.sh,v
Working file: config.tacc.sh
head: 1.10
branch:
locks: strict
access list:
symbolic names:
    start: 1.1.1.1
    TACC: 1.1.1
keyword substitution: kv
total revisions: 11;    selected revisions: 11
description:
-----
revision 1.10
date: 2007/03/30 21:23:44;  author: karl;  state: Exp;  lines: +6 -6
** Production Native IB Config - in production on 3/29/2007
-----
revision 1.9
date: 2007/03/28 22:46:55;  author: karl;  state: Exp;  lines: +2 -1
** change default stripe count from 4 to 8.
-----
```

Checking Status

```
lslogin1$ cvs status somefile
```

- Gives you
 - the repo version number
 - the local version number
 - state of the file
 - any current “sticky” tags
- Changes nothing

```
[root@lonestar2 config]# cvs status config.tacc.sh
karl@cvs.tacc.utexas.edu's password:

=====
File: config.tacc.sh      Status: Locally Modified

Working revision:      1.10
Repository revision: 1.10    /cvs/adv/system/admin/lustre/lonestar2/config.tacc.sh,v
Sticky Tag:            (none)
Sticky Date:           (none)
Sticky Options:        (none)
```


Statuses

- Up-to-date
- Locally Modified
- Locally Added
- Locally Removed
- Needs Checkout
- Needs Patch
- Needs Merge
- Unresolved Conflict (two files w/ the same name)
- File had conflicts on merge
- Unknown

Resolving Conflicts

- Open the file that is listed as having a conflict
- Search for the line that starts with “<<<<<< thisfile”
 - lists the local file name after the “<”s
 - following lines are the local version of the code
- A line that starts with “=====” marks the transition
 - the repo code follows
- A line with lots of “>>>>>> x.y” ends this conflict section
 - lists the repo version number after the “>”s

Resolving Conflicts

- Decide what code actually needs to appear in this section
- Remove the “<”, “=”, and “>” lines
- Rerun update and commit
- CAVEAT CVS will accept anything that’s in the file once you change the date of the file
 - conflict markers might be legit in your code, so be careful
 - it will warn you if you leave the conflict markers in, but will do the commit anyway

Removing Files

- If you remove a file, and then do an update, CVS will get you a fresh copy from the repo
- To remove a file from the repo, you must
 - remove the file locally
 - **then** ask CVS to remove it with the remove command
 - then commit the removal

```
lslogin1$ rm somefile; cvs remove somefile
```

```
lslogin1$ rm somefile; cvs rm somefile
```

- The file isn't actually removed, you can get it back by asking for an older version of the repo
 - CVS remembers the name, so if you add a file back, it pulls in the old history

Showing Differences

```
lslogin1$ cvs diff somefile
```

- Shows the difference between the local copy and the repo version
 - differences are shown in the format used by *diff(1)*
 - many *diff(1)* options can be applied to *cvs diff* to affect the formatting

```
lslogin1$ cvs diff -r tag somefile
```

- Shows the difference between the local copy and some other version specified by the *tag*

```
[root@lonestar2 config]# cvs diff config.tacc.sh
karl@cvs.tacc.utexas.edu's password:
Index: config.tacc.sh
=====
RCS file: /cvs/adv/system/admin/lustre/lonestar2/config.tacc.sh,v
retrieving revision 1.10
diff -r1.10 config.tacc.sh
51a52,54
> ${LMC} -o $config --add net --node ls2-mds --nid ls2-mds --nettype $LOCALNET
> ${LMC} -m $config --add net --node ls2-mds1 --nid ls2-mds1 --nettype $LOCALNET
>
60a64
> ${LMC} -m $config --add net --node data-0-$i --nid data-0-$i --nettype $LOCALNET
```

Revision Numbering and Tags

- CVS numbers every version of every file you commit
 - each file gets its own revision number history
 - numbering starts at 1.1
 - usually goes 1.1, 1.2, 1.3, ...
 - ... but other things are possible
 - generally advisable not to change CVS's numbering strategy
- You can use `cvс ci -r x.y` to specify the revision number
 - but I wouldn't unless you know what you're doing
 - `-r` can be used with other commands to get access to a particular file's revision (like before with `cvс diff`)

Tags

```
lslogin1$ cvs tag sometag
```

- Applies the symbolic tag *sometag* to every file
 - modifies the repo directly
 - associates revision numbers with *sometag*
 - uses the current revision number
 - be careful if there are locally modified files
 - CVS won't warn you
 - the repo version is the thing that gets tagged
 - cannot use '\$,.,;,@' in your tag names
 - I use: vX_Y for version X.Y
 - i.e. if this is version 3.4 of my code, I tag it with v3_4
- Symbolic tags can be used as arguments to *-r*

Using Tags

```
lslogin1$ cvs co -r sometag -d  
someotherdir somemodule
```

- Checks out a fresh copy of an older version of the module tagged with *sometag*
- This tag is “sticky”
 - CVS commands will all operate under this tag
 - *cvs up -A* will clear all sticky tags and merge the local versions with the most current repo versions

When to Commit?

- Committing too often leads may leave the repo in a state where the current version doesn't compile
- Committing too infrequently means that collaborators are wait for your important changes, bugfixes, etc. to show up
 - makes conflicts much more likely
- Common policies
 - committed files must compile and link
 - committed files must pass some test
- Come to some agreement with your collaborators about the state of the repo

Cute Tricks in Your Code

- You can keep track of the CVS revision automatically in all your source code files
- Just add `"$Id: $"` into a comment in your file
- This gets expanded when you check in/out the relevant file to include the revision, last update date, and the user making the changes
- For example:

```
/* $Id: lsf_showq.c,v 1.26 2007/10/02 18:00:52 karl Exp $ */
```

Cute Tricks in Your Code

- You can also include a record of the log history if you want
- Add “\$Log: \$” into the desired location
- Note that you only get log entries from the point at which you add the CVS Log entry
- Here is an example from a shell script

```
# $Log: mvapich-intel.spec,v $  
# Revision 1.1 2007/06/22 09:35:06 karl  
# ** Initial standard build of mvapich for Intel  
#
```

Subversion

- Subversion is a Functional superset of CVS (so if you learn cvs, you can also function in subversion)
 - Directory versioning (rename and moves)
 - Atomic Commits
 - File meta-data
 - True client-server model
 - Cross-platform, open-source

Subversion Architecture

- Each working copy has a .svn directory
 - *Similar to the CVS's CVS directory*
 - *Stores metadata about each file*
- Local or Network Repository
- Network access over HTTP or SSH
- Encrypted authentication
 - *Cleartext password stored in ~/.subversion*
- Fine-grained authorization
- Command line client is svn

CVS vs. Subversion

- Most CVS commands exist in SVN
 - *Checkout, commit, update, status...*
- Arguments position matters in CVS

```
$ cvs -d /cvsrootupdate -d
```
- Not so in SVN

```
$ svn log -r 123 foo.c
$ svn log foo.c -r 123
```

Revisions

- Revision numbers are applied to an object to identify a unique version of that object.
- CVS
 - Revision numbers are per file.
 - No connection between two files with the same revision number.
 - A commit only modifies the version number of the files that were modified.

foo.c rev 1.2 and bar.c rev 1.10.

- After commit of bar.c:

foo.c rev 1.2 and bar.c rev 1.11.

Revisions (cont)

- Revision numbers are applied to an object to identify a unique version of that object.
- SVN
 - *Revision numbers are global across the whole repository.*
 - *Snapshot in time of the whole repository.*
 - *A commit modifies the version number of all the files.*

foo.c rev 25 and bar.c rev 25.

- *After commit of bar.c:*

foo.c rev 26 and bar.c rev 26.

Subtle Note: foo.c rev 25 and 26 are identical.

References

- CVS FAQ:
http://ximbiot.com/cvs/wiki/index.php?title=CVS_FAQ
- CVS Manual
<http://ximbiot.com/cvs/manual/>