

CS395T: Introduction to Scientific and Technical Computing Debugging

Instructors:

Dr. Karl W. Schulz, TACC

Dr. Victor Eijkhout, TACC



THE UNIVERSITY OF TEXAS AT AUSTIN
Texas Advanced Computing Center

Debugging Scientific Applications

- Motivation for developing good debugging skills:
 - Unless you are from a new planet, you will introduce bugs at some point in your code
 - Even if you use community applications written by others, they will introduce bugs
 - And yes, commercial applications have bugs too
- Extra problems:
 - As scientific researchers, we cannot simply concern ourselves with bugs that prevent the application from running
 - We actually care deeply about the accuracy and repeatability of the result (*eg. negative density values in a flow code are probably bad*)
 - Stability is a concern (*eg. an iterative based solver that never converges*)
- The addition of strong debugging skills to your toolbox will greatly enhance your efficiency and add confidence to the numerical results



Defensive Programming Tips

- One of the best defenses against runtime bugs is to use basic defensive programming techniques:
 - Check **all** function return codes for errors
 - Check **all** input values controlling program execution to ensure they are within acceptable ranges (even those from flat text files in which you know there could not possibly be an error)
 - Echo all physical control parameters to a location that you will look at routinely (eg. *stdout*). Better yet, save all the parameters necessary to repeat an analysis in your solution files (*remember the metadata options in netCDF and HDF?*)
 - In addition to monitoring for obvious floating-point problems (eg. *divide-by-zero*), check for non-physical results in your simulations (*eg. supersonic velocities predicted in a low-speed aerodynamic simulation*)



Defensive Programming Tips

- Additional suggestions:
 - Maintain test cases for regression testing - is there an analytic test case you can benchmark against?
 - Use version control systems (CVS, Subversion, etc)
 - Maintain a clean, modular structure with documented interfaces: goto's, long-jumps, clever/obscure macros, etc. are dangerous over the long haul
 - Why not include some comments? Your colleagues will thank you and it just might save your dissertation when you are revisiting a tricky piece of code after a year or two
 - Strong error checking is the mark of a sage programmer and will give you more confidence in your numerical results

Defensive Programming

- Q: Isn't checking all the error codes a waste of time?
- A: It is substantially less wasteful than long debugging sessions which could have been avoided by simple error checks
- Useful error checks indicate you know what you are doing - at an **absolute minimum**, please check your memory allocations:

```
p = (float
*)calloc(nnodes+1,sizeof(float));
if(p == NULL){
    printf("Allocation error for p!\n");
    exit(1);
}
```

C

```
allocate(buf(na), stat = ierror)
if(ierror > 0) then
    print*, 'ERROR: Unable to allocate array buf'
    stop
endif
```

Fortran

Defensive Programming

- An easy defensive strategy for Fortran programmers is to use **IMPLICIT NONE** in all your routines (and specifically typecast all used variables). Avoids any undesired conversions
- Be sure to initialize all variables and arrays that require it (don't count on the architecture/OS to do this for you)
- During the testing and validation phase, make use of available compiler options to debugging options to trap:
 - Intel Fortran Examples:
 - check all** - enable runtime checks for out-of-bounds array subscripts, uninitialized variables, etc
 - warn all** - display all relevant warning messages
 - warn errors** - tells the compiler to change all warning-level messages into error-level messages
 - fpe0** - tells the compiler to abort when any floating point exceptions occur
 - GCC flags to display all warnings and catch errors:
 - Wall, -Wextra, -Wshadow, -Wunreachable-code**

Consult your compiler documentation for available runtime checks

Defensive Programming Example

- Consider the following example code ([problem.c](#)):

```
int main()
{
    int a, b;
    int x1, x2;

    if (a = b)
        printf("%d\n", x1);
    return 0;
}
```

- `> gcc -o problem problem.c`
`> ./problem`

1074729080

Use the -Wall function to help find errors at compile time

Defensive Programming Example

Now, use the “-Wall” option to have the compiler point out possible trouble spots

```
> gcc -Wall -o problem problem.c
```

```
problem.c: In function main:
```

```
problem.c:6: warning: suggest parentheses around assignment  
used as truth value
```

```
problem.c:7: warning: implicit declaration of function printf  
problem.c:7: warning: incompatible implicit declaration of  
built-in function printf
```

```
problem.c:4: warning: unused variable x2
```

*Warnings along with line
numbers are provided*

```
1      int main()  
2      {  
3          int a, b;  
4          int x1, x2;  
5  
6          if (a = b)  
7              printf("%d\n",  
x1);  
8          return 0;  
9      }
```


Defensive Programming

- Provide one or more levels of instrumentation in your code for debugging (eg. *a debug mode and a verbose debug mode*)
- C programmers should take advantage of the `assert` macro to ensure values fall within appropriate ranges

```
> gcc -o macro macro.c
> ./macro
> macro: macro.c:12: main: Assertion `n <= 100' failed.
Abort
> gcc -DNDEBUG -o macro macro.c
> ./macro
>
```

```
#include <stdio.h>
#include <assert.h>

int main()
{
    int n;
    float x[100];
    n = 1000;

    /* Assert that n <= 100 */

    assert ( n <= 100 );

    return 0;
}
```




Defensive Programming

- Q: What is the equivalent of assert in Fortran?
- You are on your own, but it is easy for programmers to improvise (some pre-processing is handy)
- This looks like mixed code, how do we compile this?

```
program main
  implicit none
  integer n
  real x(100)
  n = 1000

#ifdef DEBUG
  if ( n > 100) then
    print*, ' Assertion (n <= 100) is false'
    print*, ' File: ', __FILE__, ' Line: ', __LINE__
    stop
  endif
#endif

stop
end
```



```
> ifort -DDEBUG -cpp macro.f
> ./a.out
Assertion (n <= 100) is false
File: macro.f Line: 10
```

Bug Identification

- Common instances in which bugs identify themselves:
 - Build errors (Makefile, preprocessor, compiler, linker)
 - Improper memory reads/writes
 - pointer errors, array bounds overruns, uninitialized memory references
 - alignment problems, exhausting memory, memory leaks
 - Misinterpretation of memory
 - Type errors, e.g. when passing parameters
 - Scope/naming errors (e.g., shadowing a global name with a local name)
 - Illegal numerical operations (divide by zero, overflow, underflow)
 - Infinite loops
 - Stack overflow
 - I/O errors
 - Logic / algorithmic errors
 - Poor performance

Bug Identification

- What are the symptoms if your application has a memory bug?
 - wrong answers derived when using values from incorrect memory space
 - application behaves differently when different levels of optimization are applied; a classic memory bug symptom is as follows:
 - you compile with full optimization (-O3 for example) and your code crashes unexpectedly
 - you disable all optimization (-O0 for example) and the code runs fine
 - adding additional print statements in the program to try and isolate the bug seems to make it disappear
 - application receives an unexpected symbol and terminates
- We need to understand what it means for our application to receive an extern signal

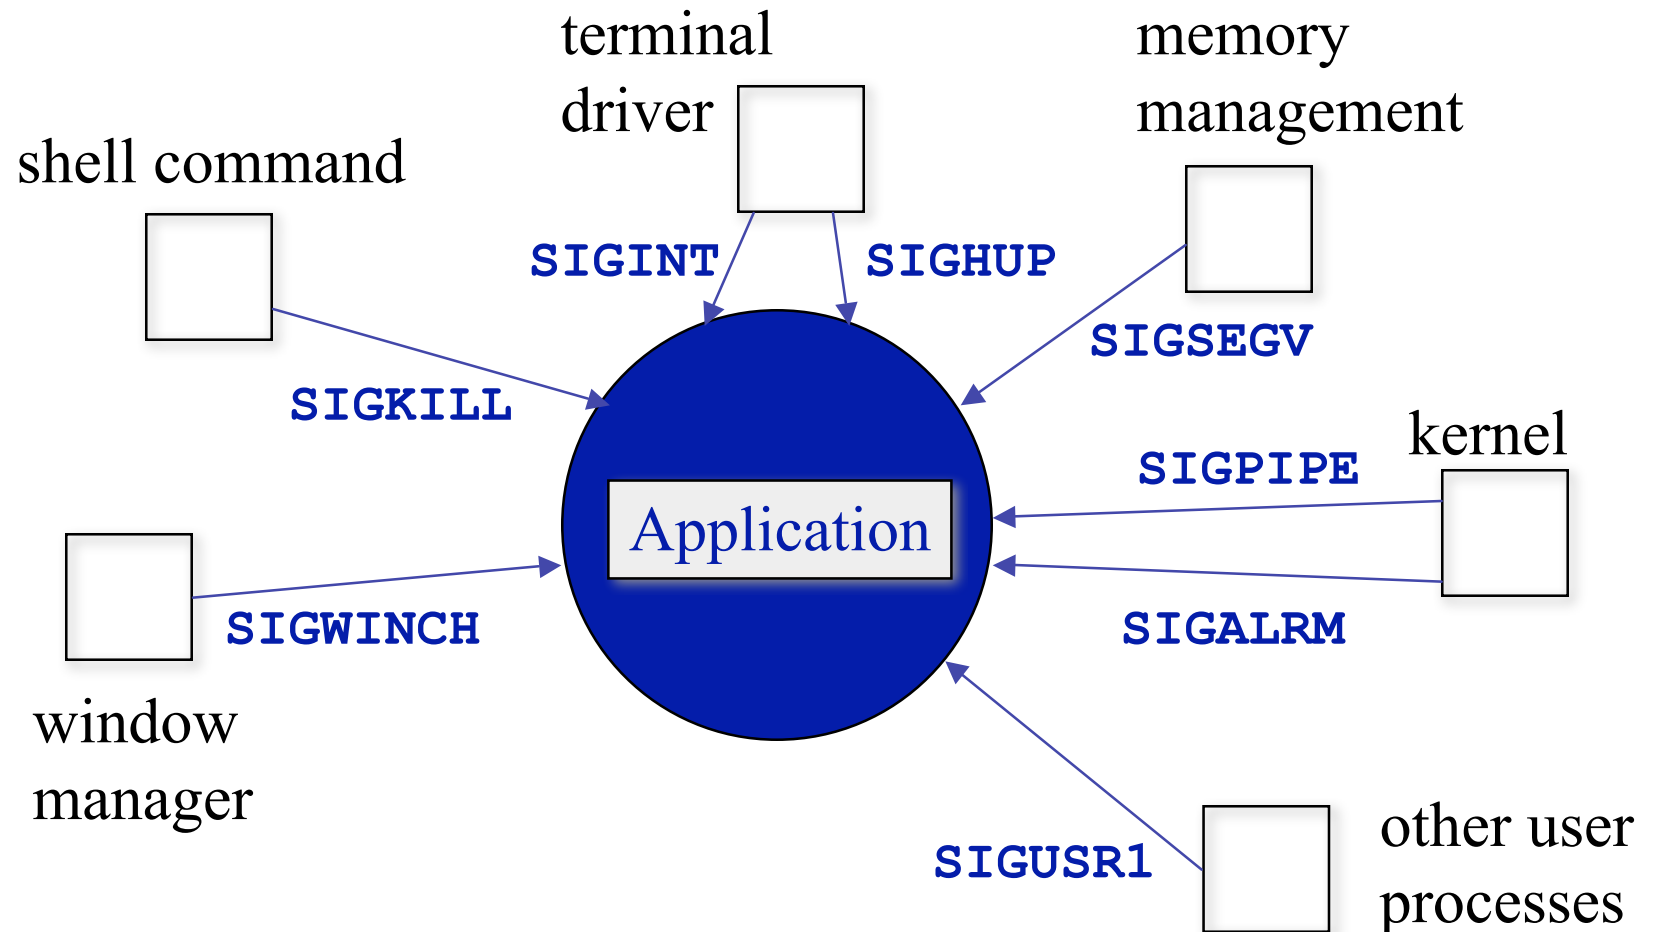
Signals

- A signal is an *asynchronous* event which is delivered to a process.
- Asynchronous means that the event can occur at any time
 - may be unrelated to the execution of the process
 - e.g. user types `ctrl-C`, or the operating system detects an error and sends a signal to your application

Common Signal Types

- | <u><i>Name</i></u> | <u><i>Description</i></u> | <u><i>Default Action</i></u> |
|--------------------|---------------------------------|------------------------------|
| SIGINT | Interrupt character typed | terminate process |
| SIGQUIT | Quit character typed (^\\) | create core image |
| SIGKILL | Kill signal | terminate process |
| SIGFPE | Floating exception | create core image |
| SIGSEGV | Invalid memory reference | create core image |
| SIGPIPE | Write on pipe but no reader | terminate process |
| SIGALRM | alarm() clock 'rings' | terminate process |
| SIGUSR1 | user-defined signal type | terminate process |
| SIGUSR2 | user-defined signal type | terminate process |
| | : | |
- See `man 7 signal` for more information

Common Signal Sources for Applications



Core Dumps

- Recall that the default action for a SIGSEGV was to create a core image and abort
- What in the world is a core image?
 - a **core dump** is a record of the raw contents of one or more regions of working memory for an application at a given time
 - commonly used to debug a program that has terminated abnormally
 - Note: if your application is “segfaulting” and you are not getting a core file, you may need to alter your shell limits. For example, in TCSH:

> `unlimit coredumpsize`
 - Main benefit of a core file is that a post-mortem analysis can be performed on an application that failed. If the symbol table was included, you can backtrace to exactly where the application when the exception occurred.
 - Note: the location of an exception may not be the original location of a bug (particularly for memory bugs)

Debugging Process

- We recognize that defensive programming can greatly reduce debugging needs, but at some point we all have to roll up our sleeves and track down a bug
- The basic steps in debugging are straightforward in principle:
 - Recognize that a bug exists
 - Isolate the source of the bug
 - Identify the cause of the bug
 - Determine a fix for the bug
 - Apply the fix and test it
- In practice, these can be difficult for particularly pesky bugs; hence, we need some more tools at our disposal (a *debugger*)

Standard Debuggers

- Command line debuggers are powerful tools to aid in diagnosing problematic applications and are available on all Unix architectures for C/C++ and Fortran
- Example debuggers:
 - Linux: gdb
 - AIX: dbx
 - SUN: dbx
 - TACC machines: DDT graphical parallel debugger
- The basic use of these debuggers is as a front-end for stepping through your application and examining variables, arrays, function returns, etc at different times during the execution
- Gives you an opportunity to investigate the dynamic runtime behavior of the application

Note: we will focus primarily on [gdb](#), but concepts and syntax are similar in [dbx](#)

Debugging Basics

- For effective debugging a couple of commands need to be mastered:
 - show program backtraces (the calling history up to the current point)
 - set breakpoints
 - display the value of individual variables
 - set new values
 - step through a program

Debugging Basics

- A **breakpoint** is a pseudo instruction that the user can insert at any place into the program during a debugging session
- Conceptually, the execution is controlled by the debugger and the debugger will interpret the breakpoints
- When execution crosses a breakpoint, the debugger will pause program execution so that you can:
 - inspect variables,
 - set or clear breakpoints, and
 - continue execution

Debugging Basics

- The notion of a **conditional breakpoint** also exists in which additional logic can be associated with the breakpoint
- When a conditional breakpoint is crossed during execution, the program will pause only if the breakpoint's break condition holds
- Example break conditions:
 - A given expression is true
 - The breakpoint has been crossed N times ("hit count") - this is very handy when you know something bad is happening on a particular iteration
 - A given expression has changed its value

Debugging Basics

- For debugging sessions, you should compile your application with extra debugging information included (eg. the symbol table)
- The symbol table maps the binary execution calls back to the original source code definitions
- To include this information, add “-g” to your compilation directives:

```
> gcc -g -o hello hello.c
```

Running GDB

- `gdb` is started directly from the shell
- You can include the name of the program to be debugged, and an optional core file:
 - `gdb`
 - `gdb a.out`
 - `gdb a.out corefile`
- `gdb` can also attach to a program that is already running; you just need to know the PID associated with the desired process

spawns a new instance of ./a.out

examines trapped state in corefile

– `gdb a.out 1134`

useful if an application seems to be slow or stuck and you want to see what it is doing currently

gdb Basics

- Common commands for gdb:
 - **run** - starts the program; if you do not set up any breakpoints the program will run until it terminates or core dumps - program command line arguments can be specified here
 - **print** - prints a variable located in the current scope
 - **next** - executes the current command, and moves to the next command in the program
 - **step** - steps through the next command. Note: if you are at a function call, and you issue **next**, then the function will execute and return. However, if you issue **step**, then you will go to the first line of that function
 - **break** - sets a break point.
 - **continue** - used to continue till next breakpoint or termination

Note: shorthand notations exist for most of these commands: eg. 'c' = continue

gdb Basics

- More commands for gdb:
 - **list** - show code listing near the current execution location
 - **delete** - delete a breakpoint
 - **condition** - make a breakpoint conditional
 - **display** - continuously display value
 - **undisplay** - remove displayed value
 - **where** - show current function stack trace
 - **help** - display help text
 - **quit** - exit gdb

gdb Basics

- Consider the following C code for subsequent examples ([hello.c](#)):

```
#include <stdio.h>
void foo();

int main()
{
    printf("inside main\n");
    foo();
    return;
}

void foo()
{
    int i, total=0;
    printf("inside foo\n");
    for(i=0;i<1000;i++)
        total += i;
}
```

Example GDB Session

```
> gcc -g -o hello hello.c
> gdb ./hello
GNU gdb Red Hat Linux (6.3.0.0-1.134.fc5rh)
Copyright 2004 Free Software Foundation, Inc.

(gdb) run
Starting program: /home/karl/cs395t/hello
inside main
inside foo

Program exited with code 0347.

(gdb) break main
Breakpoint 1 at 0x8048384: file hello.c, line 5.

(gdb) run
Starting program: /home/karl/cs395t/hello

Breakpoint 1, main () at hello.c:5
5      {
(gdb) where
#0  main () at hello.c:5
```

Example GDB Session (continued)

```
(gdb) break foo
```

```
Breakpoint 2 at 0x80483b5: file hello.c, line 13.
```

```
(gdb) cont
```

```
Continuing.
```

```
inside main
```

```
Breakpoint 2, foo () at hello.c:13
```

```
13      int i, total=0;
```

```
(gdb) list
```

```
8          return;
```

```
9      }
```

```
10
```

```
11      void foo() {
```

```
12      {
```

```
13          int i, total=0;
```

```
14          printf("inside foo\n");
```

```
15          for(i=0;i<1000;i++) {
```

```
16              total += i;
```

```
17      }
```

Example GDB Session (continued)

```
(gdb) cont
```

```
Continuing.
```

```
inside foo
```

```
Program exited with code 0347.
```

```
(gdb) delete breakpoints 1 2
```

```
(gdb) break hello.c:16
```

```
Breakpoint 3 at 0x80483d1: file hello.c, line 16.
```

```
(gdb) condition 3 i==401
```

```
(gdb) run
```

```
Starting program: /home/karl/cs395t/hello
```

```
inside main
```

```
inside foo
```

```
Breakpoint 3, foo () at hello.c:16
```

```
16          total += i;
```

```
(gdb) print i
```

```
$1 = 401
```

```
(gdb) where
```

```
#0  foo () at hello.c:16
```

```
#1  0x080483a6 in main () at hello.c:7
```

gdb and Emacs

- Emacs has a lot of functionality to support software development and debugging:
 - Supports compilations (eg. can access your Makefile)
 - Can running gdb directly
 - Can integrate with your revision control system (eg. CVS)
- To illustrate, consider the *hello.c* example and a trivial Makefile

```
all: hello

hello: hello.c
    gcc -g -o hello hello.c
```

gdb and Emacs

- To work within emacs, load up the src file:
`> emacs hello.c`
- Then, to compile, issue `M-x compile` which will run `make -k` by default
- emacs will create a new window for you and show you the result of your compilation
- To run the program within the debugger, issue `M-x gdb` and provide the name of your executable (eg. `hello`)
- emacs will then provide you an interface to gdb and show you the current execution location in your src code

gdb and Emacs

```
File Edit Options Buffers Tools Gud Complete In/Out Signals Help
#include <stdio.h>
void foo();

int main()
=>
    printf("inside main\n");
    foo();
    return;
}

--u:---F1 hello.c (C Abbrev)--L5--Top-----

(gdb) break main
Breakpoint 1 at 0x8048384: file hello.c, line 5.
(gdb) run
Starting program: /home/karl/cs395t/hello
Reading symbols from shared object read from target memory...done.
Loaded system supplied DSO at 0x24f000

Breakpoint 1, main () at hello.c:5
(gdb) 
--uu:**-F1 *gud-hello* (Debugger:run)--L43--Bot-----
```


gdb and Emacs

```
File Edit Options Buffers Tools Gud Complete In/Out Signals Help
```

```
void foo()  
{  
  int i, total=0;  
  printf("inside foo\n");  
  for(i=0;i<1000;i++)  
=>  total += i;  
}
```

```
--u:---F1 hello.c (C Abbrev)--L16--Bot-----
```

```
Reading symbols from shared object read from target memory...done.  
Loaded system supplied DSO at 0x24f000
```

```
Breakpoint 1, main () at hello.c:5
```

```
(gdb) step 8
```

```
inside main
```

```
inside foo
```

```
(gdb) print i
```

```
$1 = 1
```

```
(gdb) █
```

```
--uu:**-F1 *gud-hello* (Debugger:run)--L48--Bot-----
```

Memory debugging in gdb

```
#include <stdio.h>
int main()
{
    double *a;
    a = 2;
    a[1] = 5.1;
    return 0;
}
```

Null pointers and illegal addresses
are easy to recognize

```
(gdb) run
Starting program:a.out
Reading symbols for shared libraries . done
```

```
Program received signal EXC_BAD_ACCESS, Could not access memory
Reason: KERN_PROTECTION_FAILURE at address: 0x0000000a
0x00001f98 in main () at segfault.c:6
(gdb) print a
$1 = (double *) 0x2
(gdb)
```

References/Acknowledgements

- Debugging with GDB:
http://sources.redhat.com/gdb/current/onlinedocs/gdb.html#SEC_Top
- Debug malloc library: <http://dmalloc.com/>
- Electric Fence
<http://directory.fsf.org/ElectricFence.html>
- Data Display Debugger (DDD):
<http://www.gnu.org/manual/ddd/>

References/Acknowledgements

- Arctic Region Supercomputing Center, Core Skills for Computational Science:
<http://people.arsc.edu/~cskills/>
- Debugging with GDB:
http://sources.redhat.com/gdb/current/online/docs/gdb.html#SEC_Top