

# CS395T: Introduction to Scientific and Technical Computing

Scientific Libraries  
BLAS, LAPACK, and FFTW

*Instructors*

Dr. Victor Eijkhout, Research Scientist, TACC  
Dr. Karl W. Schulz, Research Associate, TACC



THE UNIVERSITY OF TEXAS AT AUSTIN  
**Texas Advanced Computing Center**

---

# Scientific Libraries

- BLAS (<http://www.netlib.org/blas/>)
  - Basic Linear Algebra Subprograms
  - basic operations on vectors and matrices
- LAPACK (<http://www.netlib.org/lapack/>)
  - Linear Algebra PACKage
  - linear systems solvers (regular and least-squares)
  - eigensolvers
  - singular value decomposition
- FFTW (<http://www.fftw.org/>)
  - Fastest Fourier Transform in the West
  - DFT (complex or real to complex)
  - DCT and DST (real to real)

# BLAS

- Three levels
  - 1 – Scalar, Vector, and Vector-Vector operations
  - 2 – Matrix-Vector operations
  - 3 – Matrix-Matrix operations
- Originally written in FORTRAN, posted to Netlib
  - uses FORTRAN array storage
- Optimized implementations now common
  - Vendors
    - Intel: MKL (<http://www.intel.com/cd/software/products/asmo-na/eng/perflib/mkl/index.htm>)
    - AMD: ACML (<http://developer.amd.com/acml.jsp>)
    - IBM: ESSL (<http://www-03.ibm.com/systems/p/software/essl.html>)
  - Third-Party
    - ATLAS (<http://math-atlas.sourceforge.net/>) automatic search
    - GotoBLAS (<http://www.tacc.utexas.edu/resources/software/#blas>)  
superman programmer approach

# LAPACK

- FORTRAN routines for linear algebra
  - uses FORTRAN array storage
- Optimized versions available
  - vendors (same as the BLAS)
    - may not implement every function
    - usually pretty complete
  - third-party
    - less coverage
    - ATLAS implements only certain routines
- Replaces
  - LINPACK and EISPACK
  - uses block operations
  - calls the BLAS where possible

# FFTW

- ANSI C code callable from C/C++ and FORTRAN
- Written by a code generator written in OCaml
- Used in
  - signal/image processing
  - spectral methods for PDEs

# Matrix Data Layout—C vs. FORTRAN

- FORTRAN arrays
  - indices start at 1
  - 2-dimensional arrays
    - implemented as contiguous storage
    - stored down the columns
      - perhaps, counter-intuitive
      - stride-1 access goes down a column (i.e. increments  $i$  first)
  - multi-dimensional arrays
    - contiguous
    - left-most argument is stride-1
- Arrays in C
  - indices start at 0
  - pointers and array indices are equivalent (usually)
  - i.e.,  $a[5]$  *means*  $*(a+5)$

# Matrix Data Layout—C vs. FORTRAN

- Multi-dimensional arrays in C/C++ are tricky
- Look like pointers of pointers
  - i.e. `a[3][4]` *can mean* `*(* (a+3) +4)`
  - Actual implementation is as contiguous data
- Can't pass a multi-dimensional array to a subroutine (unlike Fortran, C99)
- Easier to emulate:
  - Instead of `a[3][4]` write `a[3*n+4]`

# Matrix Data Layout—C vs. FORTRAN

- To have C and FORTRAN interoperate
  - must use the FORTRAN ordering
    - C is flexible on ordering (pointers to pointers)
    - FORTRAN isn't
    - so the FORTRAN ordering “wins”
  - must use contiguous storage
  - use 1-d “polynomial” indexing in C



# Matrix Data Layout

*foo.f:*

```
program foo
  real*8 a(2,2)
  a(1,1)=1
  a(1,2)=2
  a(2,1)=3
  a(2,2)=4
  call prnt(a,2)
  call addrpoly(a,2)
  call arrnot(a,2)
  stop
end
```

*bar.c:*

```
#include <stdio.h>
void prnt_(double *a, int *n)
{
  int i=0;
  for(i=0;i<(*n)*(*n);++i)
  {
    printf("%g\n",a[i]);
  }
}
...
```

localhost\$ ./foo

```
1
3
2
4
...
```

$$\mathbf{A} = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$$

# Matrix Data Layout

- Two options in C
  - use an addressing polynomial
    - use the CPP to help ( )
    - compiler may not optimize the loops very well
      - may have to do some hand-tuning
  - use multi-dimensional array notation
    - Needs pointers
    - have to transpose the arguments
    - compiler may or may not understand what you're doing
    - No Blas3 available
- Real solution: try to use library for low level primitives
- In C++, check out Boost (<http://www.boost.org/>) for some options
  - extensions of the STL to numerical work
  - usually quite optimizable by the compiler
  - won't beat the vendor or 3<sup>rd</sup>-party BLAS in most cases

# Matrix Data Layout

*bar.c:*

```
...  
#define A(i,j) a[(i)+(j)*(*n)]  
void addrpoly_(double*a,int *n)  
{  
    int i, j;  
    for (i=0; i<(*n); ++i)  
    {  
        for (j=0; j<(*n); ++j)  
            printf("%g ",A(i,j));  
        printf("\n");  
    }  
}  
...
```

```
localhost$ ./foo  
...  
1 2  
3 4  
...  
localhost$
```

# Matrix Data Layout

*bar.c:*

```
....  
void arnnot_(double *a, int *n)  
{  
    double **b;  
    int i, j;  
    for (i=0;i<(*n);++i)  
        b[i]=(a+(i*(*n)));  
    for (i=0;i<(*n);++i)  
    {  
        for (j=0;j<(*n);++j)  
            printf("%g ",b[j][i]);  
        printf("\n");  
    }  
}
```

localhost\$ ./foo

....

1 2

3 4

localhost\$

# BLAS

- Dense matrix operations
  - also includes symmetric (complex Hermitian) , triangular, and banded, but the storage is weird
- Level 1—vector-vector
  - copy
  - dot product
  - scale and shift
  - rotations (plane, Givens, etc.)
- Level 2—matrix-vector
  - matrix-vector product (plus a scale and shift)
  - rank-1 updates
- Level 3—matrix-matrix
  - matrix-matrix product (plus a scale and shift)

# BLAS Nomenclature

- Function names are prefixed with a letter corresponding to the precision and type of the input:

`xDOT`

`xAXPY`

`xGEMV`

`xGEMM`

- Prefixes

- S, single-precision real
  - `real*4` and `float`
  - ...on most architectures
- D, double-precision real (`real*8` and `double`)
- C, single-precision complex (`complex*8` and ...)
- Z, double-precision complex (`complex*16` and ...)

# Complex Data in C/C++

- C99 has a `complex` type qualifier

```
#include <complex.h>
```

```
...
```

```
double complex a[100];
```

- defines `I` to represent  $(-1.0)^{1/2}$

```
double complex b = 1.0 + 3.0*I;
```

- `complex.h` defines

- `creal()`, `cimag()`, `carg()`, `cabs()`,
- `csin()`, `ccos()`, `ctan()`, `cexp()`, `cpow()`,
- ...

- C++ has `std::complex<T>`
- Both are bitwise-equivalent to contiguous storage, i.e.

```
typedef double my_complex[2];
```

# Dot Product

$$\rho = \mathbf{x} \cdot \mathbf{y} = \mathbf{x}^T \mathbf{y}$$

```
res = ddot( n, x, incx, y, incy )
```

- `x, y, real*8` arrays length `n` (`integer*8`)
- `incx, incy` increment for `x, y`
- `res, real*8` return value
- `ddot, real*8` function (don't forget to declare this!)



# Scale a Vector

$$\mathbf{y} = a\mathbf{x} + \mathbf{y}$$

call `daxpy( n, a, x, incx, y, incy )`

- `n, x, y, incx, incy` same as for the dot product
- `a, real*8` the scaling factor

# Matrix-Vector Product

$$\mathbf{y} = \alpha \mathbf{A}\mathbf{x} + \beta \mathbf{y}$$

```
call dgemv( trans, m, n, alpha, a, lda, x, incx,  
           beta, y, incy )
```

- `m,n, integer*8` matrix dimensions
- `a, real*8` m-by-n 2-d array
- `x, real*8` vector length `m`
- `y, real*8` vector length `n`
- `incx, incy` as before
- `alpha, beta, real*8` scale factors
- `lda, integer*8` the leading dimension of `a`, usually `m`
- `trans, character*1` whether to do normal ('N'), conjugate transposed ('C'), or transposed ('T') mode

# LDA

Matrix of  $m \times n$ : element  $(i,j)$  is  
at address  $(j-1)*m+i-1$

lda parameter necessary if  
stored size different from logical  
size

```
real A(20,12)
```

```
call dgemv(... A, 7,20,4..)
```

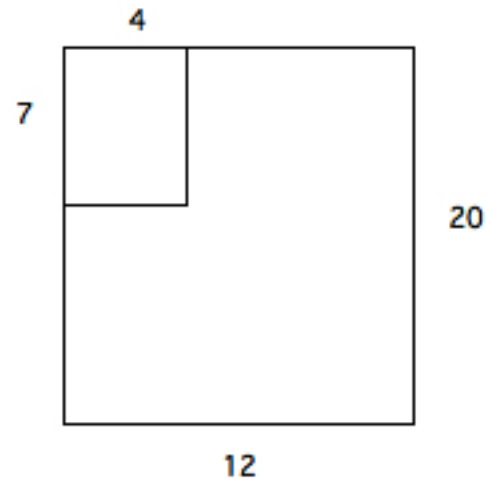
```
subroutine dgemv(...A,m,lda,n...)
```

```
real A(lda,n)
```

```
integer row,col
```

```
do row=1,n
```

```
do col=1,m
```



# Matrix-Matrix Product

$$\mathbf{C} = \alpha \mathbf{AB} + \beta \mathbf{C}$$

```
call dgemm(transa, transb, m, n, k, alpha, a, lda,  
          b, ldb, beta, c, ldc)
```

- `m, n, a, alpha, beta, lda` as before
- `k, integer*8` outermost matrix dimension
- `b, real*8 n-by-k` 2-d array
- `c, real*8 m-by-k` 2-d array
- `ldb, integer*8` leading dimension of `b` (usually `n`)
- `ldc, integer*8` leading dimension of `c` (usually `m`)
- `transa, transb` transpose mode for `a` and `b`

# LAPACK

- Can solve linear systems
  - LU factorization
  - full, symmetric, or banded storage
- Also solves least-squares and eigenproblems
  - QR factorization
  - LQ factorization
  - SVD
  - symmetric/non-symmetric eigenvalue problems
- Naming convention identical to the BLAS
- Calls the BLAS where possible

# Factor a Linear System

$$\mathbf{A} = \mathbf{PLU}$$

call `dgetrf(m, n, a, lda, ipiv, info)`

- Factorizes  $A$ , in-place, into  $L$  and  $U$  (doesn't store the unit diagonal of  $L$ )
- `m, n, integer*8` matrix dimensions
- `a, real*8` `m-by-n` 2-d array
- `lda, integer*8` leading dimension of `a` (usually `m`)
- `ipiv, integer*8` vector, length `n`, stores pivoting information
- `info, integer*8`
  - 0 for success
  - $-i$  for error in the  $i^{\text{th}}$  parameter
  - $i$  for  $U_{ii}=0$  (application will cause division by zero)

# Solve a Factored System

$$\mathbf{Ax} = \mathbf{PLUx} = \mathbf{b}$$

```
call dgetrs(trans, n, nrhs, a, lda, ipiv, b, ldb,  
info)
```

- Solves a pre-factored (as from dgetrf),  $n$ -by- $n$  linear system and overwrites  $b$  with the result
- $n, a, lda, ipiv, trans$  as before
- $b$ ,  $\text{real}^*8$   $n$ -by- $nrhs$  2-d array
- $nrhs$ ,  $\text{integer}^*8$  number of right hand sides to solve (usually 1)
- $ldb$ ,  $\text{integer}^*8$  leading dimension of  $b$  (usually  $n$ )
- $info$ ,  $\text{integer}^*8$ 
  - 0 for success
  - $-i$  for error in the  $i^{\text{th}}$  parameter

# Factor and Solve a System

$$\mathbf{Ax} = \mathbf{b}$$

```
call dgesv(n, nrhs, a, lda, ipiv, b, ldb,  
info)
```

- Solves a  $n$ -by- $n$  linear system and overwrites  $a$  and  $b$  with the factorization and the result
- All parameters as for `dgetrs`, except
- `info`, `integer*8`
  - 0 for succes
  - $-i$  for error in the  $i^{\text{th}}$  parameter
  - $i$  for  $U_{ii}=0$  (application will cause division by zero)



# A C Example

```
#include <stdlib.h>
#include <stdio.h>
#define A(i,j) a[(i) + (j)*n]
void dgetrf_();void dgetrs_();
void dgesv_();
int main()
{
    int i,j,n=3,one=1,info;
    double *a=malloc(n*n*sizeof(double));
    double *b=malloc(n*sizeof(double));
    int *ipiv=malloc(n*sizeof(int));
    char trans[1]='N';
    A(0,0)=6.;  A(0,1)=-2.;  A(0,2)=2.;
    A(1,0)=12.; A(1,1)=-8.;  A(1,2)=6.;
    A(2,0)=3.;  A(2,1)=-13.; A(2,2)=3.;
    b[0]=16.;   b[1]=26.;    b[2]=-19.;
```

```
    printf("\n\n");
    /* dgetrf_(&n,&n,a,&n,ipiv,&info);
       dgetrs_(trans,&n,&one,a,
               &n,ipiv,b,&n,&info); */
    dgesv_(&n,&one,a,&n,ipiv,b,&n,&info);
    for(i=0;i<n;++i)
    {
        printf("| ");
        for(j=0;j<n;++j)
        {
            printf("%8.5g ",A(i,j));
        }
        printf("|      | %8.5g\n",b[i]);
    }
    return(0);
}
```

# In FORTRAN

```
program bar
  integer*8 i,j,n,one
  parameter (n=3)
  parameter (one=1)
  integer*8 ipiv(n),info
  real*8 a(n,n)
  real*8 b(n)
  a(1,1)=6.
  a(1,2)=-2.
  a(1,3)=2.
  a(2,1)=12.
  a(2,2)=-8.
  a(2,3)=6.
  a(3,1)=3.
  a(3,2)=-13.
  a(3,3)=3.

  b(1)=16.
  b(2)=26.
  b(3)=-19.
  call dgesv(n,one,a,n,
             ipiv,b,n,info)

  do i=1,n
    do j=1,n
      write(*,'(F10.5$)') a(i,j)
    enddo
    write(*,'(10X,F8.5)') b(i)
  enddo
  stop
end
```

# Compiling and Linking

- On Lonestar
- Use the MKL
- Set up the environment

```
lslogin1$ module load mkl
```

- Compile and link

- C

```
icc -o foo foo.c -L$TACC_MKL_LIB -lmkl -  
lmkl_lapack -lpthread -Wl,-rpath,  
$TACC_MKL_LIB
```

- FORTRAN

```
ifort -o bar bar.f -I$TACC_MKL_INC -  
L$TACC_MKL_LIB -lmkl -lmkl_lapack -  
lpthread -Wl,-rpath,$TACC_MKL_LIB
```

# Fourier Transform

- Recall

$$F(u(x)) = \hat{u}(\xi) = \int_{-\infty}^{\infty} u(x) e^{-2\pi i x \xi} dx \quad (\text{forward})$$

$$F^{-1}(\hat{u}(\xi)) = u(x) = \int_{-\infty}^{\infty} \hat{u}(\xi) e^{2\pi i x \xi} d\xi \quad (\text{backward or inverse})$$

- Useful properties

- linear,  $F(af+bg) = aF(f) + bF(g)$
- converts differentiation, order  $n$ , to multiplication by  $(2\pi i \xi)^n$
- and integration into division by the same factor
- turns convolution,  $f^*g$ , into multiplication,  $F(f)F(g)$

# FFTW

- Computes Discrete Fourier Transforms

- forward

$$Y_k = \sum_{j=0}^{n-1} X_j e^{-2\pi j k \sqrt{-1}/n}$$

- inverse

$$Z_k = \sum_{j=0}^{n-1} Y_j e^{2\pi j k \sqrt{-1}/n}$$

- Forward followed by inverse leaves  $Z=n*X$
    - Positive frequencies stored in the first half of  $Z$  and negative in the second half
- Does 2- and 3-D DFTs, too
- Can do them in parallel (in an older version)

# Interface

- FFTW3
  - FFTW2 still around (and available on Lonestar)
  - 2 → 3 changed the API
- FORTRAN
  - uses the `complex*16` intrinsic datatype
- C/C++
  - defines `fftw_complex`  
`typedef double fftw_complex[2];`
  - can use C99's `complex`  
`double complex *x;`
  - can probably use C99's `complex` in C++ codes for most compilers or `std::complex` with some care

# Interface

- Define your data
- Make a *plan*
- Execute the plan
  - perhaps several times with new data
- Destroy the plan
- Clean up your data

# Data

```
double complex *in, *out;  
fftw_plan p;  
in = (double complex*)  
    fftw_malloc(sizeof(double complex)  
        * n);  
out = (double complex*)  
    fftw_malloc(sizeof(double complex)  
        * n);
```

- `fftw_malloc()` guarantees proper data alignment for the fastest code



# Plans

```
p = fftw_plan_dft_1d(n, in, out,  
    FFTW_FORWARD, FFTW_ESTIMATE);  
fftw_execute(p);  
fftw_destroy_plan(p);
```

- Sets up a 1-D Forward DFT
- Estimates the fastest method (FFTW\_MEASURE for best method)
- Executes the plan
- Can modify `in` and re-execute without destroying as long as `n` doesn't change

# Data Cleanup

```
fftw_free(in);
```

```
fftw_free(out);
```

- Analogous to malloc()/free()
- Should always call when you're done with the data

# 1-D Example

- Random time series in  $[0,1]$
- Transform
- Remove DC-offset
- Inverse Transform
- Low-pass filter
  - keep the first 20 modes
  - reuse the inverse transform plan

# C Code

```
#include <stdio.h>
#include <stdlib.h>
#include <complex.h>
#include "fftw3.h"
int main(int argc, char* argv[])
{
    int n=atoi(argv[1]);
    int i, cutoff = atoi(argv[2]);
    double *x=malloc(n*sizeof(double));
    for(i=0; i<n; ++i) x[i]=((double) rand())/((double)RAND_MAX+1.0);
    printf("x=[");
    for(i=0; i<n; ++i) printf("%15.7g ",x[i]);
    printf("];\n\n");
    double complex *in, *out;
    fftw_plan p;
    in = (double complex*) fftw_malloc(sizeof(double complex) * n);
    out = (double complex*) fftw_malloc(sizeof(double complex) * n);
    for(i=0; i<n; ++i) in[i]=(x[i]+0.*I);
    p = fftw_plan_dft_1d(n, in, out, FFTW_FORWARD, FFTW_ESTIMATE);
    fftw_execute(p);
    fftw_destroy_plan(p);
}
```

# C Code

```
printf("y=["); for(i=0; i<n; ++i)
    printf("%15.7g + %15.7gi ",
        1./((double) n)*creal(out[i]), 1./((double) n)*cimag(out[i]));
printf("];\n\n");
out[0]=0.+0.*I;
p = fftw_plan_dft_1d(n, out, in, FFTW_BACKWARD, FFTW_ESTIMATE);
fftw_execute(p);
printf("z=["); for(i=0; i<n; ++i)
    printf("%15.7g + %15.7gi ",
        1./((double) n)*creal(in[i]), 1./((double) n)*cimag(in[i]));
printf("];\n\n");
for(i=1; i<n/2+1; ++i)
    if (i > cutoff)
        out[i]=out[n-i]=0.+0.*I;
printf("r=["); for(i=0; i<n; ++i)
    printf("%15.7g + %15.7gi ",
        1./((double) n)*creal(out[i]), 1./((double) n)*cimag(out[i]));
printf("];\n\n");
fftw_execute(p);
printf("q=["); for(i=0; i<n; ++i)
    printf("%15.7g + %15.7gi ",
        1./((double) n)*creal(in[i]), 1./((double) n)*cimag(in[i]));
printf("];\n\n");
fftw_destroy_plan(p); fftw_free(in); fftw_free(out);
}
```

# Compiling and Linking

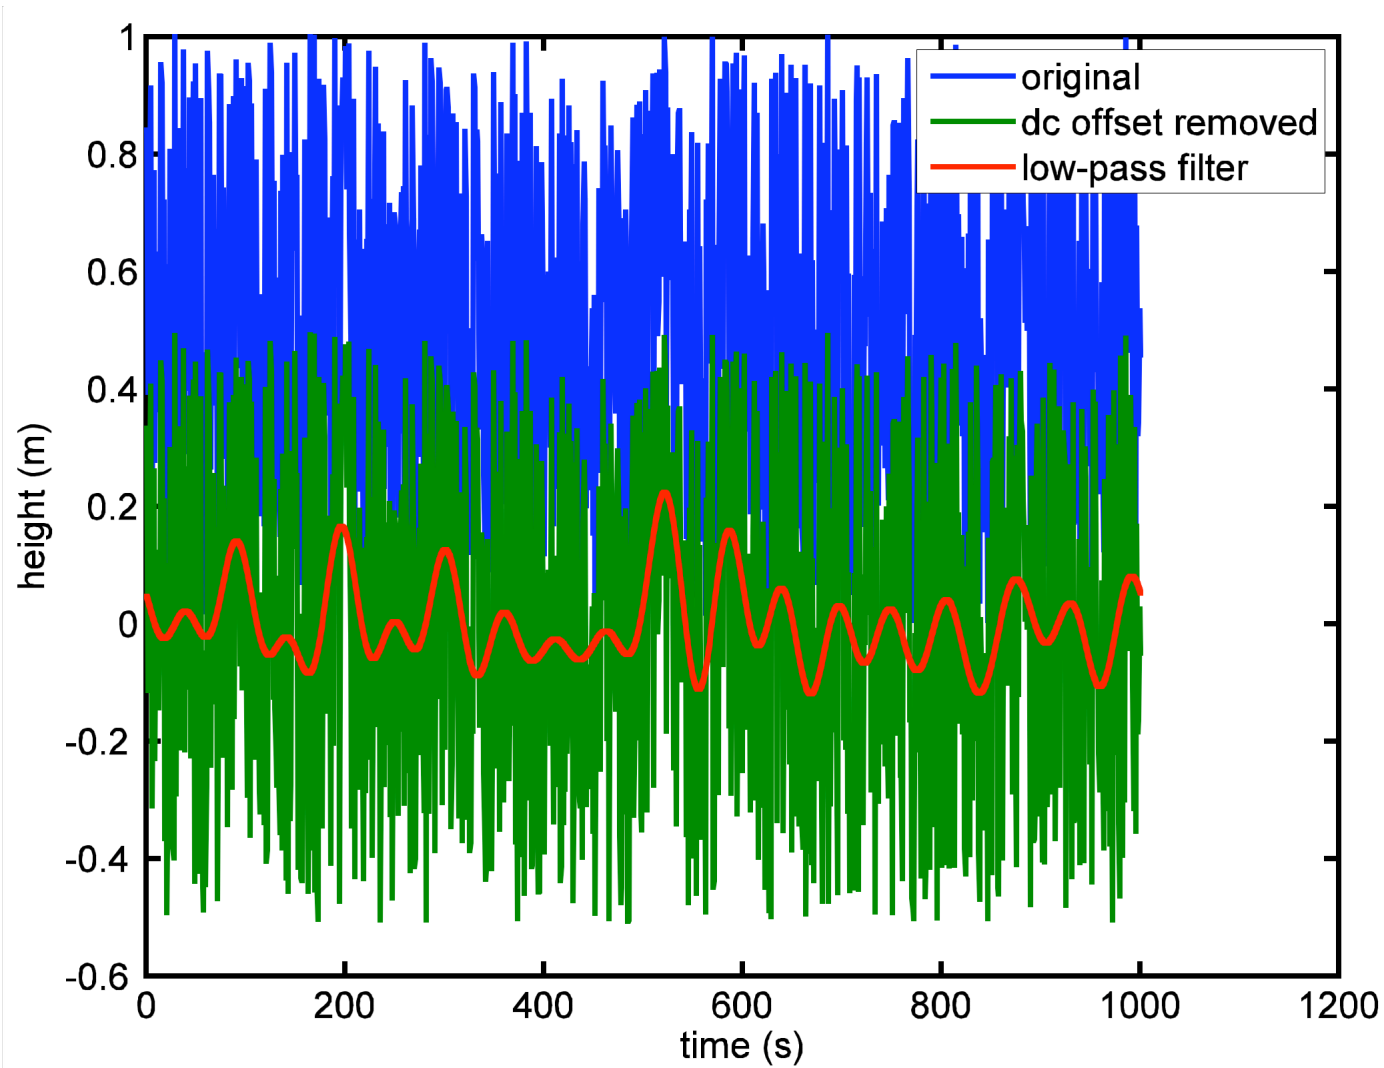
- On Lonestar
- Set up the environment

```
lslogin1$ module load fftw/3.1.1
```

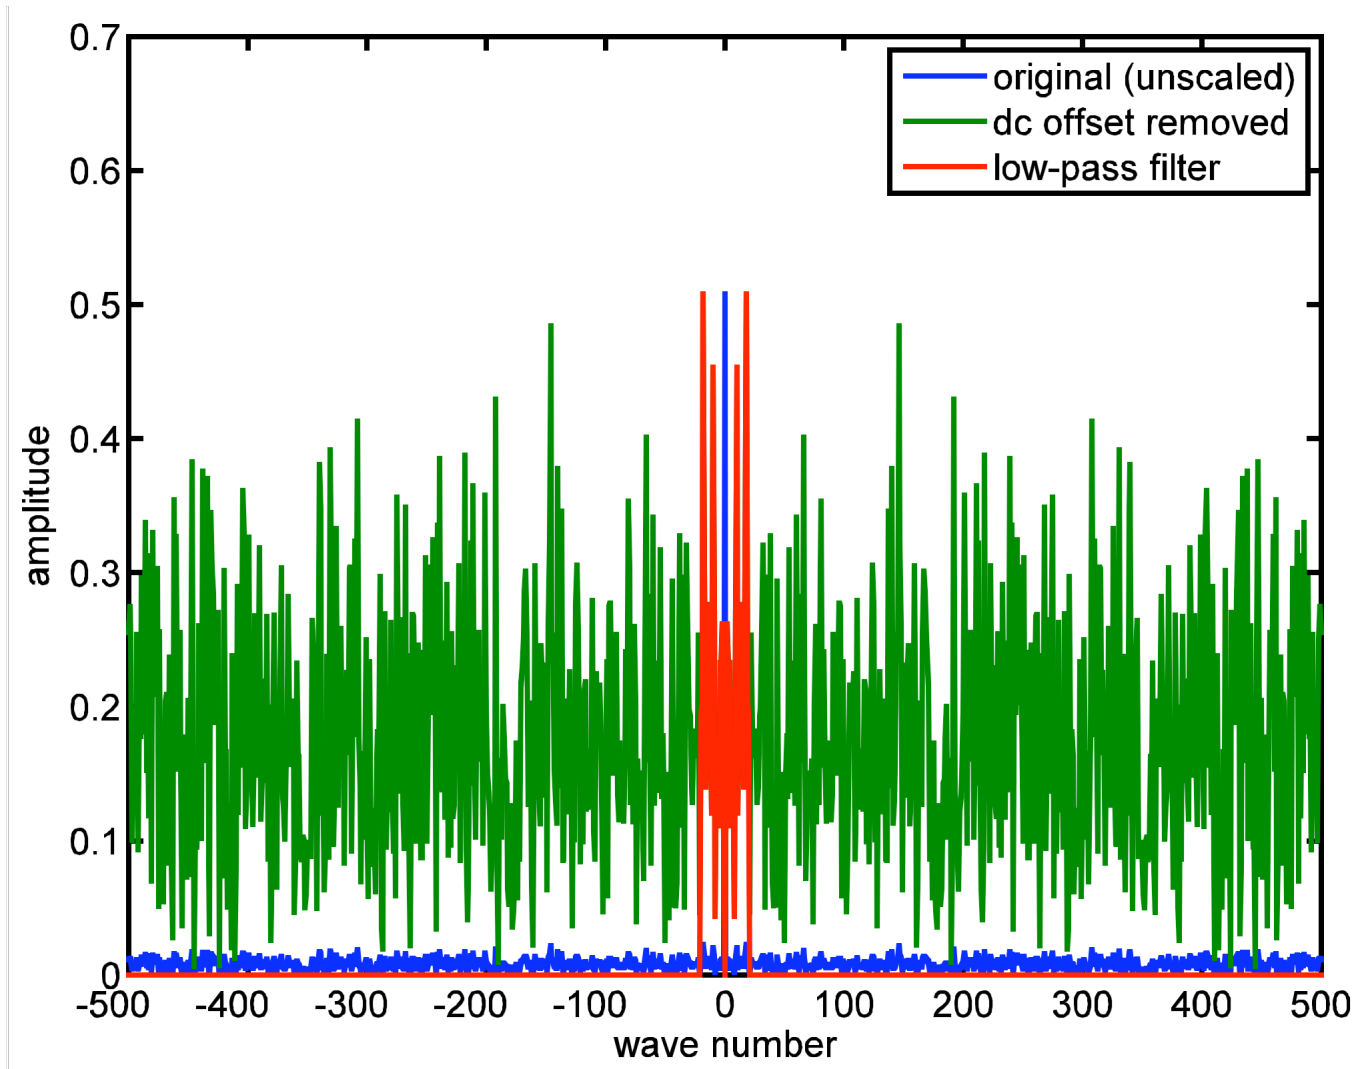
- Compile and link

```
icc -o rand_data rand_data.c -  
    L$TACC_FFTW3_LIB -lfftw3 -lm -  
    I$TACC_FFTW3_INC
```

# Results



# Results





# Results

